

▼ **Name: Soumya Ranjan Nayak**

PRN: 23070243063

Assignment: Convolution and Pooling Operations on Input Image

```
In [ ]: import kagglehub
puneet6060_intel_image_classification_path = kagglehub.dataset_download('puneet6060/intel-image-classification')
print('Data source import complete.')
```

```
In [ ]: # import library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import os
import glob as gb
import cv2
import keras
from tensorflow.keras.models import Sequential, Model
```

```
In [ ]: trainpath = '/kaggle/input/intel-image-classification/seg_train/seg_train'
testpath = '/kaggle/input/intel-image-classification/seg_test/seg_test'
predpath = '/kaggle/input/intel-image-classification/seg_pred/seg_pred'
```

```
In [ ]: IMAGE_SIZE = (228, 228)

BATCH_SIZE = 32
```

```
In [ ]: train_ds = tf.keras.utils.image_dataset_from_directory(
    trainpath,
    seed=123,
    image_size=IMAGE_SIZE,
    batch_size=BATCH_SIZE)
```

Found 14034 files belonging to 6 classes.

```
In [ ]: test_ds = tf.keras.utils.image_dataset_from_directory(
    testpath,
    seed=123,
    image_size=IMAGE_SIZE,
    batch_size=BATCH_SIZE)
```

Found 3000 files belonging to 6 classes.

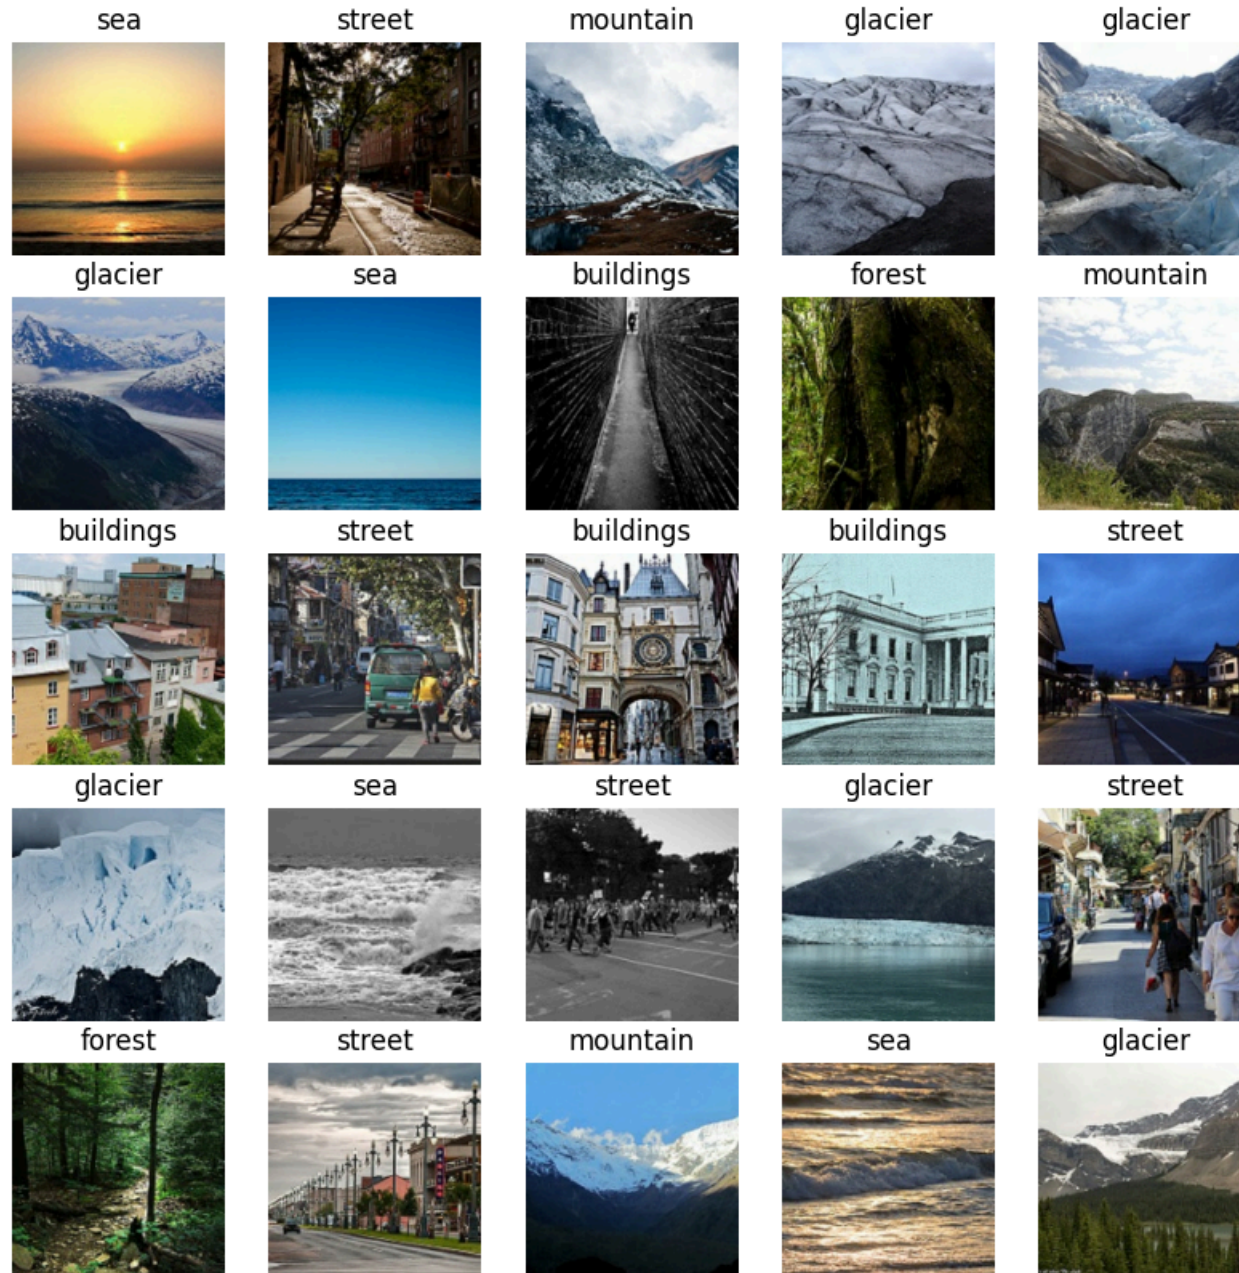
```
In [ ]: class_names = train_ds.class_names  
print(class_names)
```

```
['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

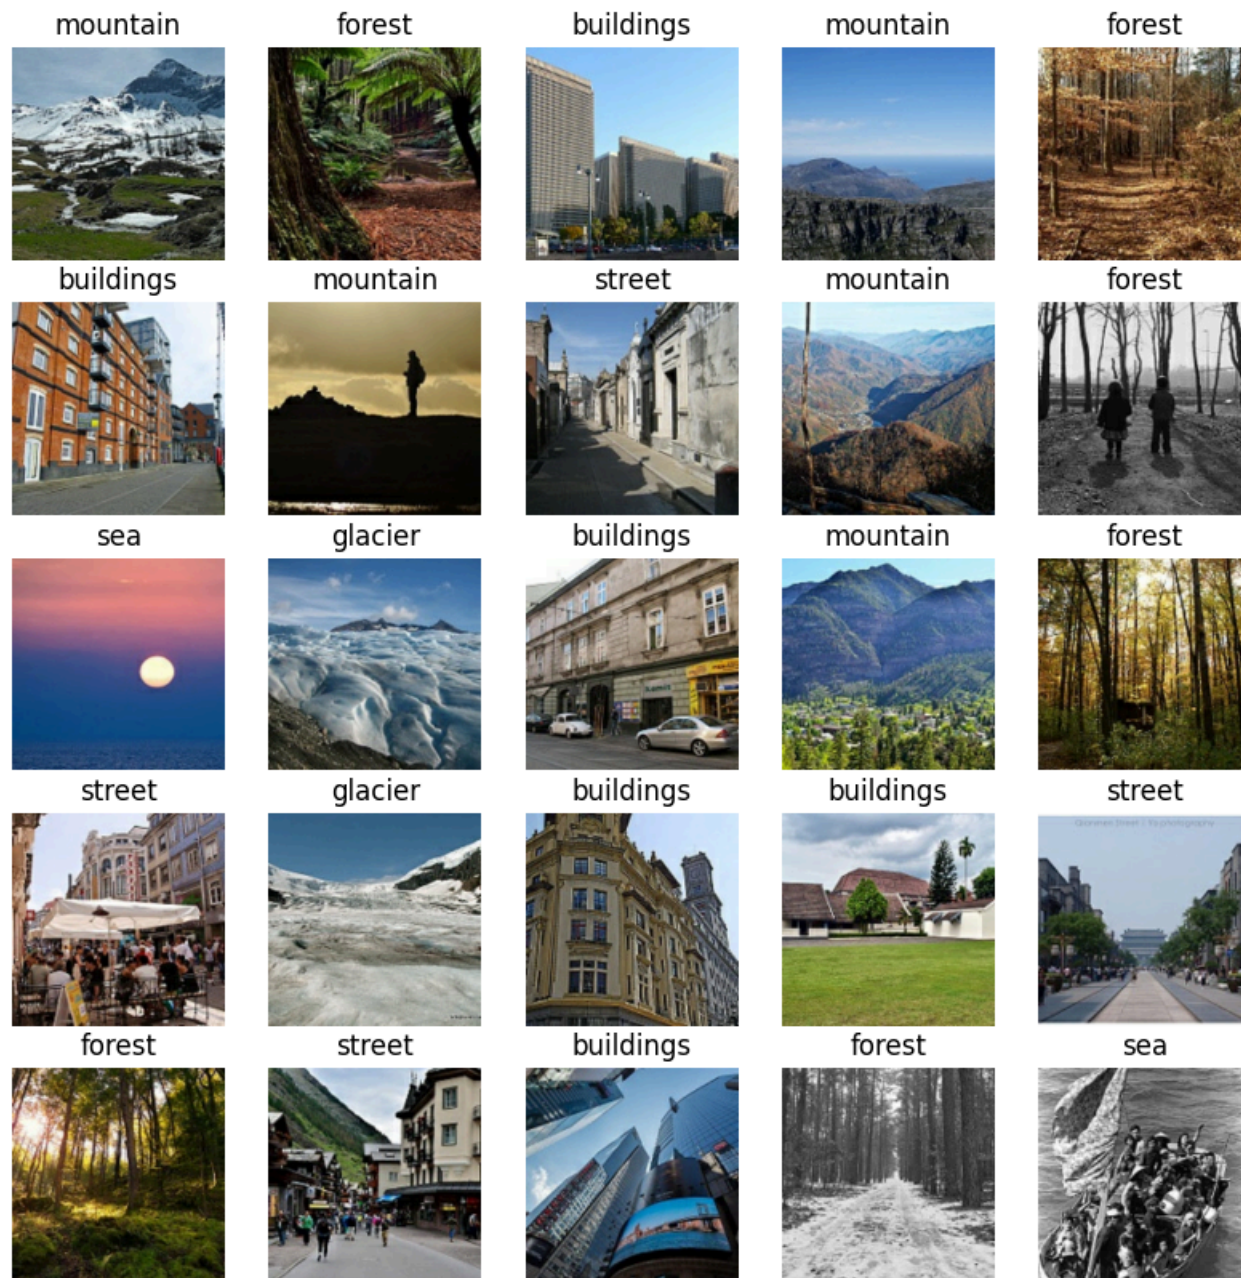
```
In [ ]: def getImagePaths(path):  
    image_names = []  
    for dirname, _, filenames in os.walk(path):  
        for filename in filenames:  
            fullpath = os.path.join(dirname, filename)  
            image_names.append(fullpath)  
    return image_names  
images_paths = getImagePaths(predpath)  
len(images_paths)
```

```
Out[13]: 7301
```

```
In [ ]: plt.figure(figsize=(10, 10))
        for images, labels in train_ds.take(1):
            for i in range(25):
                ax = plt.subplot(5, 5, i + 1)
                plt.imshow(images[i].numpy().astype("uint8"))
                plt.title(class_names[labels[i]])
                plt.axis("off")
```



```
In [ ]: plt.figure(figsize=(10, 10))
for images, labels in test_ds.take(1):
    for i in range(25):
        ax = plt.subplot(5, 5, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```




```
In [ ]: import tensorflow.keras.models as Models
```

```
In [ ]: model = Models.Sequential()  
model.add(tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(228,228,3)))  
model.add(tf.keras.layers.MaxPooling2D(2,2))  
model.add(tf.keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu'))  
model.add(tf.keras.layers.MaxPooling2D(2,2))  
model.add(tf.keras.layers.Conv2D(64, kernel_size=(3, 3), activation='relu'))  
model.add(tf.keras.layers.MaxPooling2D(2,2))  
model.add(tf.keras.layers.Conv2D(64, kernel_size=(3, 3), activation='relu'))  
model.add(tf.keras.layers.MaxPooling2D(2,2))  
model.add(tf.keras.layers.Conv2D(64, kernel_size=(3, 3), activation='relu'))  
model.add(tf.keras.layers.MaxPooling2D(2,2))  
model.add(tf.keras.layers.Flatten())  
model.add(tf.keras.layers.Dense(1024, activation='relu'))  
model.add(tf.keras.layers.Dropout(0.2))  
model.add(tf.keras.layers.Dense(128, activation='relu'))  
model.add(tf.keras.layers.Dropout(0.2))  
model.add(tf.keras.layers.Dense(len(class_names), activation='softmax'))
```

```
In [ ]: model.summary()
```

Model: "sequential_1"


| Layer (type) | Output Shape | Param # |
|-------------------------------------|----------------------|---------|
| ===== | | |
| conv2d_5 (Conv2D) | (None, 226, 226, 32) | 896 |
| max_pooling2d_5 (MaxPooling2D) | (None, 113, 113, 32) | 0 |
| conv2d_6 (Conv2D) | (None, 111, 111, 32) | 9248 |
| max_pooling2d_6 (MaxPooling2D) | (None, 55, 55, 32) | 0 |
| conv2d_7 (Conv2D) | (None, 53, 53, 64) | 18496 |
| max_pooling2d_7 (MaxPooling2D) | (None, 26, 26, 64) | 0 |
| conv2d_8 (Conv2D) | (None, 24, 24, 64) | 36928 |
| max_pooling2d_8 (MaxPooling2D) | (None, 12, 12, 64) | 0 |
| conv2d_9 (Conv2D) | (None, 10, 10, 64) | 36928 |
| max_pooling2d_9 (MaxPooling2D) | (None, 5, 5, 64) | 0 |
| flatten_1 (Flatten) | (None, 1600) | 0 |
| dense_2 (Dense) | (None, 1024) | 1639424 |
| dropout_2 (Dropout) | (None, 1024) | 0 |
| dense_3 (Dense) | (None, 128) | 131200 |
| dropout_3 (Dropout) | (None, 128) | 0 |
| dense_4 (Dense) | (None, 6) | 774 |
| ===== | | |
| Total params: 1873894 (7.15 MB) | | |
| Trainable params: 1873894 (7.15 MB) | | |
| Non-trainable params: 0 (0.00 Byte) | | |

Visualize Model

```
In [ ]: from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt

# Visualize the ANN Model
plot_model(model, to_file="ann_model.png", show_shapes=True, show_layer_names=True)
```

Out[23]:



```
In [ ]: from tensorflow.keras.optimizers import Adam
model.compile(
    optimizer = Adam(learning_rate = 0.001),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    #loss = "categorical_crossentropy",
    metrics = ["accuracy"])
```

```
In [ ]: from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping

earlystopping = EarlyStopping(monitor='val_loss',
                              patience=5,
                              verbose=1,
                              mode='min'
                              )

checkpointer = ModelCheckpoint(filepath='bestvalue.keras', verbose=0, save_best_only=True)
callback_list = [checkpointer, earlystopping]
```

```
In [ ]: history = model.fit(train_ds,
    validation_data=test_ds,
    epochs=40,
    callbacks=callback_list
)
```

Epoch 1/40

/opt/conda/lib/python3.10/site-packages/keras/src/backend.py:5727: UserWarning: "`sparse_categorical_crossentropy` received `from_logits=True`, but the `output` argument was produced by a Softmax activation and thus does not represent logits. Was this intended?

output, from_logits = _get_logits(

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1731841900.673195 101 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

439/439 [=====] - 42s 79ms/step - loss: 1.2959 - accuracy: 0.5594 - val_loss: 0.9318 - val_accuracy: 0.6540

Epoch 2/40

439/439 [=====] - 19s 42ms/step - loss: 0.7888 - accuracy: 0.7098 - val_loss: 0.7498 - val_accuracy: 0.7017

Epoch 3/40

439/439 [=====] - 19s 43ms/step - loss: 0.6503 - accuracy: 0.7698 - val_loss: 0.6610 - val_accuracy: 0.7700

Epoch 4/40

439/439 [=====] - 19s 42ms/step - loss: 0.5669 - accuracy: 0.7990 - val_loss: 0.5911 - val_accuracy: 0.8030

Epoch 5/40

439/439 [=====] - 19s 42ms/step - loss: 0.4990 - accuracy: 0.8244 - val_loss: 0.5630 - val_accuracy: 0.8070

Epoch 6/40

439/439 [=====] - 18s 41ms/step - loss: 0.4485 - accuracy: 0.8364 - val_loss: 0.8302 - val_accuracy: 0.7403

Epoch 7/40

439/439 [=====] - 18s 41ms/step - loss: 0.4163 - accuracy: 0.8495 - val_loss: 0.6385 - val_accuracy: 0.7960

Epoch 8/40

439/439 [=====] - 18s 41ms/step - loss: 0.3804 - accuracy: 0.8617 - val_loss: 0.6208 - val_accuracy: 0.8087

Epoch 9/40

439/439 [=====] - 18s 41ms/step - loss: 0.3420 - accuracy: 0.8772 - val_loss: 0.6726 - val_accuracy: 0.8123

Epoch 10/40

439/439 [=====] - 18s 41ms/step - loss: 0.3042 - accuracy: 0.8928 - val_loss: 0.6881 - val_accuracy: 0.7947

Epoch 10: early stopping

▼ Accuracy

```
In [ ]: # Evaluate the model on the test set
loss1, accuracy1 = model.evaluate(test_ds, verbose=0)
print(f"Test Accuracy: {accuracy1:.2f}")
```

Test Accuracy: 0.79

▼ Save Model

```
In [ ]: # Save the model to a file
model.save("model.h5")
print("Model saved as 'model.h5'")
```

Model saved as 'model.h5'

/opt/conda/lib/python3.10/site-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.
saving_api.save_model(

▼ Plotting training and validation loss over epochs

```

In [ ]: # Assuming `history` contains the training history
history_df1 = pd.DataFrame(history.history)



# Create a figure with 1 row and 2 columns for side-by-side plots
fig, axes = plt.subplots(1, 2, figsize=(12, 4), sharex=True)

# Plot accuracy on the first subplot
axes[0].plot(history_df1['accuracy'], "#6daa9f", label='Training Accuracy')
axes[0].plot(history_df1['val_accuracy'], "#774571", label='Validation Accuracy')
axes[0].set_title('Training and Validation Accuracy')
axes[0].set_xlabel('Epochs')
axes[0].set_ylabel('Accuracy')
axes[0].legend(loc='best')

# Plot Loss on the second subplot
axes[1].plot(history_df1['loss'], "#ffad5a", label='Training Loss')
axes[1].plot(history_df1['val_loss'], "#af4b64", label='Validation Loss')
axes[1].set_title('Training and Validation Loss')
axes[1].set_xlabel('Epochs')
axes[1].set_ylabel('Loss')
axes[1].legend(loc='best')

# Adjust layout for better spacing
plt.tight_layout()
plt.show()

```



▼ User Input

```
In [ ]: def predict_image(filename, model):  
        img_ = image.load_img(filename, target_size=(228, 228))  
        img_array = image.img_to_array(img_)  
        img_processed = np.expand_dims(img_array, axis=0)  
        img_processed /= 255.  
  
        prediction = model.predict(img_processed)  
  
        index = np.argmax(prediction)  
  
        plt.title("Prediction - {}".format(str(classes[index]).title()), size=18)  
        plt.imshow(img_array)
```

```
In [ ]: from tensorflow.keras.preprocessing import image  
        predict_image('/kaggle/input/intel-image-classification/seg_pred/seg_pred/1003.jpg', model)
```

Prediction - Sea

```
In [ ]:
```